



UNC CHARLOTTE

*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING

# Introduction to ML

## Lecture 2: Linear Regression

Hamed Tabkhi

Department of Electrical and Computer Engineering,  
University of North Carolina Charlotte (UNCC)

[htabkhiv@uncc.edu](mailto:htabkhiv@uncc.edu)



UNC CHARLOTTE

# Linear regression

---

- In regression problem, the task is:  
approximate a mapping function ( $h$ ) from input variables ( $x$ ) to continuous output variables
- In this lecture, we work on linear regression models

# Linear regression

---

Linear regression is a simple approach to supervised learning. It assumes that the dependence of  $Y$  on  $X_1, X_2, \dots, X_p$  is linear.

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots \beta_k x_k$$

  
linear  
combination

$\beta$  unknown model parameters are estimated from the data

# Linear regression

---

- A linear regression model follows a very particular form. In statistics, a regression model is linear when all terms in the model are one of the following:
  - The constant
  - A **parameter** multiplied by an **independent variable** (IV)

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots \beta_k x_k$$

# Quiz (1)

---

- Linear or non-linear regression model

$$\theta_1 * X^{\theta_2}$$


non-linear

$$\theta_1 * \cos(X + \theta_4) + \theta_2 * \cos(2 * X + \theta_4) + \theta_3$$

non-linear

$$Y = b_0 + b_1X_1 + b_2X_1^2$$

Non-linear



Q: Are the true regression functions always linear ?

# Linear regression

---

- Why should we assume that relationships between variables are linear?
  - Because linear relationships are the simplest non-trivial relationships that can be imagined (hence the easiest to work with), and.....
  - Because the "true" relationships between our variables are often at least approximately linear over the range of values that are of interest to us, and...
  - Even if they're not, we can often transform the variables in such a way as to linearize the relationships.


# Linear regression

- Let's start with an example

Training set

( <b>x</b> <b>ft<sup>2</sup></b> , \$ <b>y</b> K)
(3883   ft <sup>2</sup> , \$432K)
(1668   ft <sup>2</sup> , \$218K)
(3577   ft <sup>2</sup> , \$366K)
(1668   ft <sup>2</sup> , \$218K)
(765   ft <sup>2</sup> , \$123K)
(3822   ft <sup>2</sup> , \$493K)
⋮

Predict house price






# Linear regression

---


- Linear regression with one variable (x)



# Linear regression

---

- Linear regression with one variable (x) – simple linear regression



Linear function


$$h(x) = \theta_0 + \theta_1 x$$

$\theta_0, \theta_1$  are parameters

# Linear regression

---

- How to choose the parameters?
- Idea: choose  $\theta_0, \theta_1$  so that  $h(x)$  is close to  $y$  for our training examples  $(x, y)$



# Linear regression

---

- Let's formalize the problem

**Goal: Obtaining a Linear Regression Model**

Input:  $x$  (the size of house)

Target:  $y$  (the price of the house),

Linear model:

$$h = \theta_0 + \theta_1 x$$

# Linear regression

- **Cost Function: How well or poorly a model (Hypothesis) explains the training data**

Sum of the squared Euclidean distance

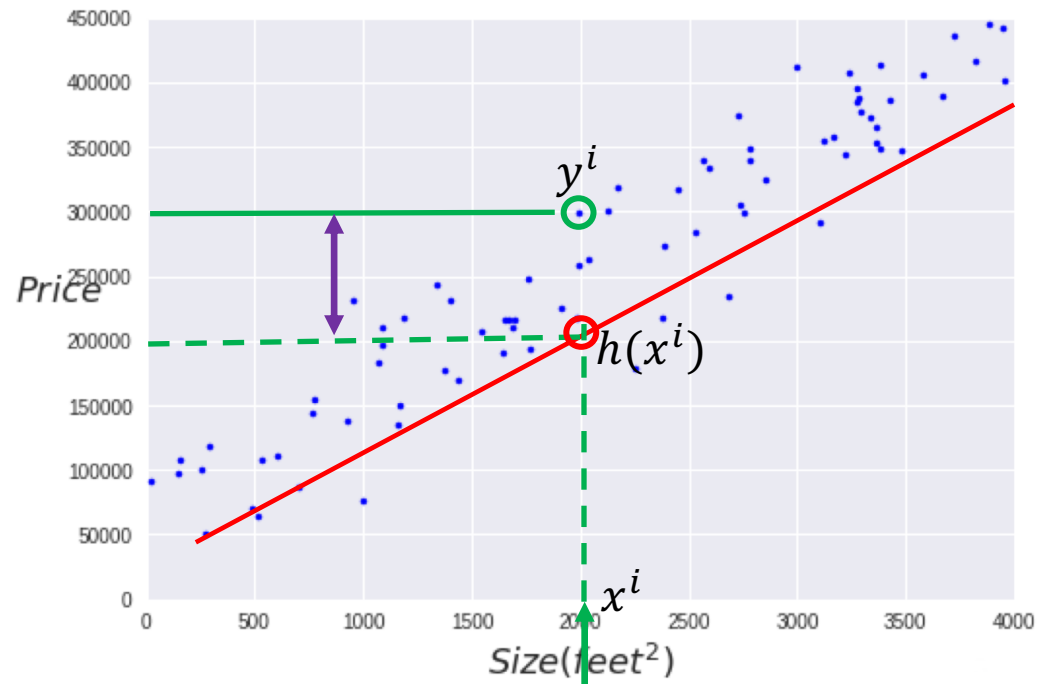
$$J(\theta_0, \theta_1) = \sum_i (h(x^i) - y^i)^2$$

$(x^i, y^i)$  is the  $i$ th training sample

$$h(x^i) = \theta_0 + \theta_1 x^i$$

$$i \in [1, m]$$

$m$ : total number of training samples



# Linear regression

---

- Problem to solve

Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters:  $\theta_0, \theta_1$

Cost Function:  $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

$m$ : total number of training samples

$\frac{1}{2}$ : mathematical convenience


Goal: minimize  $J(\theta_0, \theta_1)$

Cost function

# Linear regression

---

Visualizing cost function  $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$



# Linear regression



- Simplified hypothesis  $h$  for easier visualization

Set  $\theta_0 = 0$

$$h_{\theta}(x) = \theta_1 x$$

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (\theta_1 x^{(i)} - y^{(i)})^2$$



a quadratic polynomial




# Gradient descent algorithm

- How to find the minimum point of cost function?
  - Gradient descent
    - Used all over machine learning for minimization
- But wait

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (\theta_1 x^{(i)} - y^{(i)})^2$$

$$\frac{dJ(\theta_1)}{d\theta_1} = 0$$




Analytical Method

# Gradient descent algorithm

---

- Start with an initial guess
- Keeping changing  $\theta_1$  a little bit to try and reduce  $J(\theta_1)$



# Gradient descent algorithm

---


Have a cost function  $J(\theta_0, \theta_1)$

Want  $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

## Outline:

- Start with some  $\theta_0, \theta_1$
- Keep changing  $\theta_0, \theta_1$  to reduce  $J(\theta_0, \theta_1)$   
until we hopefully end up at a minimum

# Gradient descent algorithm



# Gradient descent algorithm

---

- A more formal definition

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

}

Annotations:

- $\theta_j$  (next): next
- $\theta_j$  (current): current
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ : **Gradient:** partial derivative
- $\alpha$  is learning rate (a positive constant)
- Note: j is an index for parameters

## Correct: Simultaneous update

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$

# Gradient descent algorithm

---

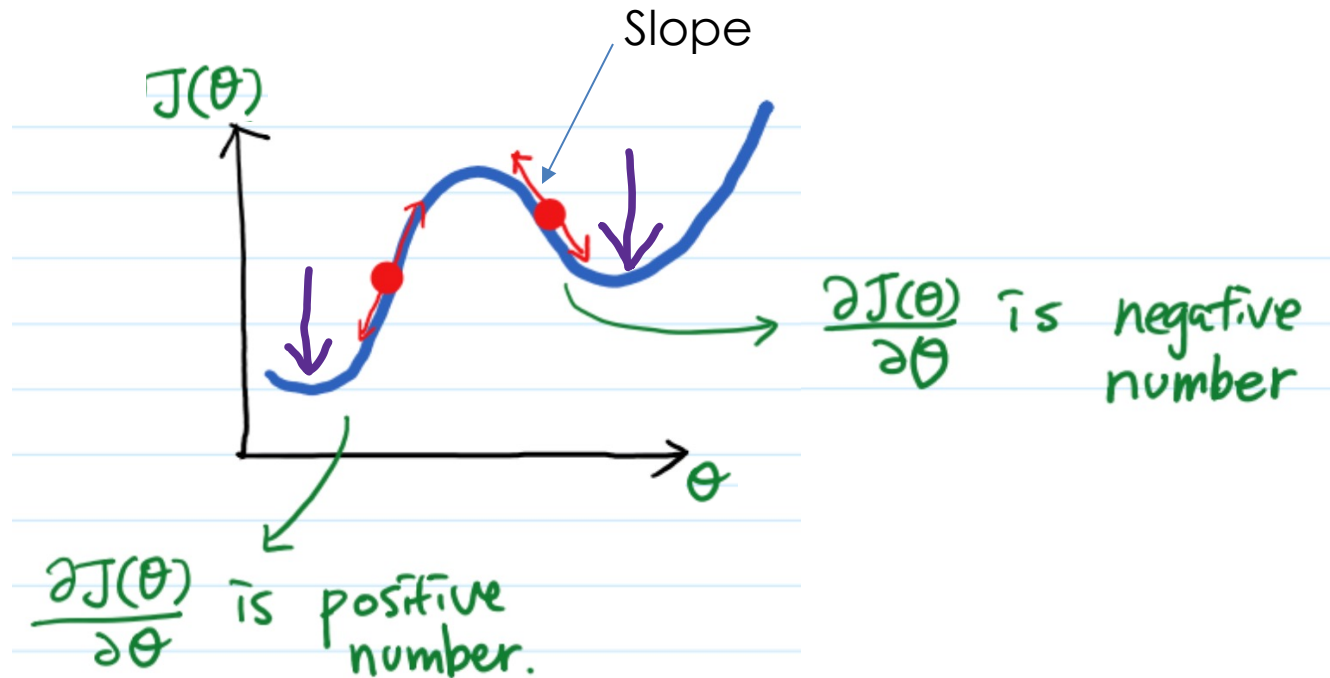
- Q1: Why “ - ” the gradient in parameter update?
- Q2: How to set a proper value for the learning rate  $\alpha$  ?

# Gradient descent algorithm

---

- Note that a gradient is a vector, so it has both of the following characteristics:
  - a direction
  - a magnitude

# Gradient descent algorithm




$$\theta = \theta - \alpha \cdot \text{slope}$$

The learning rate  $\alpha$  is a positive constant



# Gradient descent algorithm



$$\theta = \theta - \alpha \cdot \text{slope}$$


1.9  $\theta_{\text{new}} = 2 - \alpha \cdot (2) \quad (\theta_{\text{old}} = 2)$

4.1  $\theta_{\text{new}} = 4 - \alpha (-2) \quad (\theta_{\text{old}} = 4)$

$$\alpha = 0.05$$

# Gradient descent algorithm


- How to choose a proper learning rate  $\alpha$  ?



$$\theta = \theta - \alpha \cdot \text{slope}$$

1.9  $\theta_{\text{new}} = 2 - \alpha \cdot (2) \quad (\theta_{\text{old}} = 2)$


4.1  $\theta_{\text{new}} = 4 - \alpha \cdot (-2) \quad (\theta_{\text{old}} = 4)$   
 $\alpha = 0.05$



Gradually approach the minimum

# Gradient descent algorithm


---



If the learning rate is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge.


# Gradient descent algorithm

Too low




A small learning rate requires many updates before reaching the minimum point

Just right



The optimal learning rate swiftly reaches the minimum point

Too high



Too large of a learning rate causes drastic updates which lead to divergent behaviors

# Gradient descent algorithm

---

- Fixed learning rate
  - determined by trial and error
  - E.g., try a few values 0.1, 0.01, 0.001, ..., i.e., parameter tuning via grid search
- To see if gradient descent is working, print out  $J(\theta)$  each iteration
  - The value should decrease at each iteration
  - If it doesn't, adjust  $\alpha$


# Gradient descent algorithm

---

- Recommended reading on learning rate

<https://medium.com/@lipeng2/cyclical-learning-rates-for-training-neural-networks-4de755927d46>

Learning rate annealing, initially large, decrease gradually



Cyclical Learning Rates for Training Neural Networks

If the learning rate is too small, then gradient descent may take a very long time to converge.

If  $\theta_0$  and  $\theta_1$  are initialized at a local minimum, then one iteration will not change their values.

Even if the learning rate  $\alpha$  is very large, every iteration of gradient descent will decrease the value of  $f(\theta_0, \theta_1)$ .

If  $\theta_0$  and  $\theta_1$  are initialized so that  $\theta_0 = \theta_1$ , then by symmetry (because we do simultaneous updates to the two parameters), after one iteration of gradient descent, we will still have  $\theta_0 = \theta_1$ .

# Linear regression with multiple variables

More input variables (features)

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)		Price (\$1000)
$x_1$	$x_2$	$x_3$	$x_4$	$x_n$	$y$
2104	5	1	45	...	460
1416	3	2	40		232
1534	3	2	30		315
852	2	1	36		178

Each row is an n-dimensional data point (sample)



# Linear regression with multiple variables

---

- Previous hypothesis or mapping function

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Now 
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convenience of notation, define  $x_0 = 1$ .

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} \in R^{n+1} \qquad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \in R^{n+1}$$



Each data point is now a  $(n+1)$ -dimensional vector

# Linear regression with multiple variables

Hypothesis:  $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

Parameters:  $\theta_0, \theta_1, \dots, \theta_n$

Cost function:  $m$ : total number of training samples

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

*i-th sample* *Corresponding label or GT value*

Repeat {

- Gradient descent  $\rightarrow \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$   
} (simultaneously update for every  $j = 0, \dots, n$ )

$i$ : index for sample

$j$ : index for parameter

# Linear regression with multiple variables

Hypothesis:  $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$

Parameters:  $\theta_0, \theta_1, \dots, \theta_n$

$m$ : total number of training samples

Cost function:

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Repeat {

$\downarrow \frac{2}{2\theta_j} J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update  $\theta_j$  for  
 $j = 0, \dots, n$ ) }

$i$ : index for sample  
 $j$ : index for parameter

# Linear regression with multiple variables

---

- Stopping criterion

- Assume convergence when  $\|\boldsymbol{\theta}_{new} - \boldsymbol{\theta}_{old}\|_2 < \epsilon$

A small constant (threshold)

- Set a maximum number of iterations

# Useful resource

---

- Gradient and partial derivatives
  - <https://www.youtube.com/watch?v=GkB4vW16QHI>
- A very detailed tutorial on Gradient Descent (Step-by-Step) with examples
  - <https://www.youtube.com/watch?v=sDv4f4s2SB8>